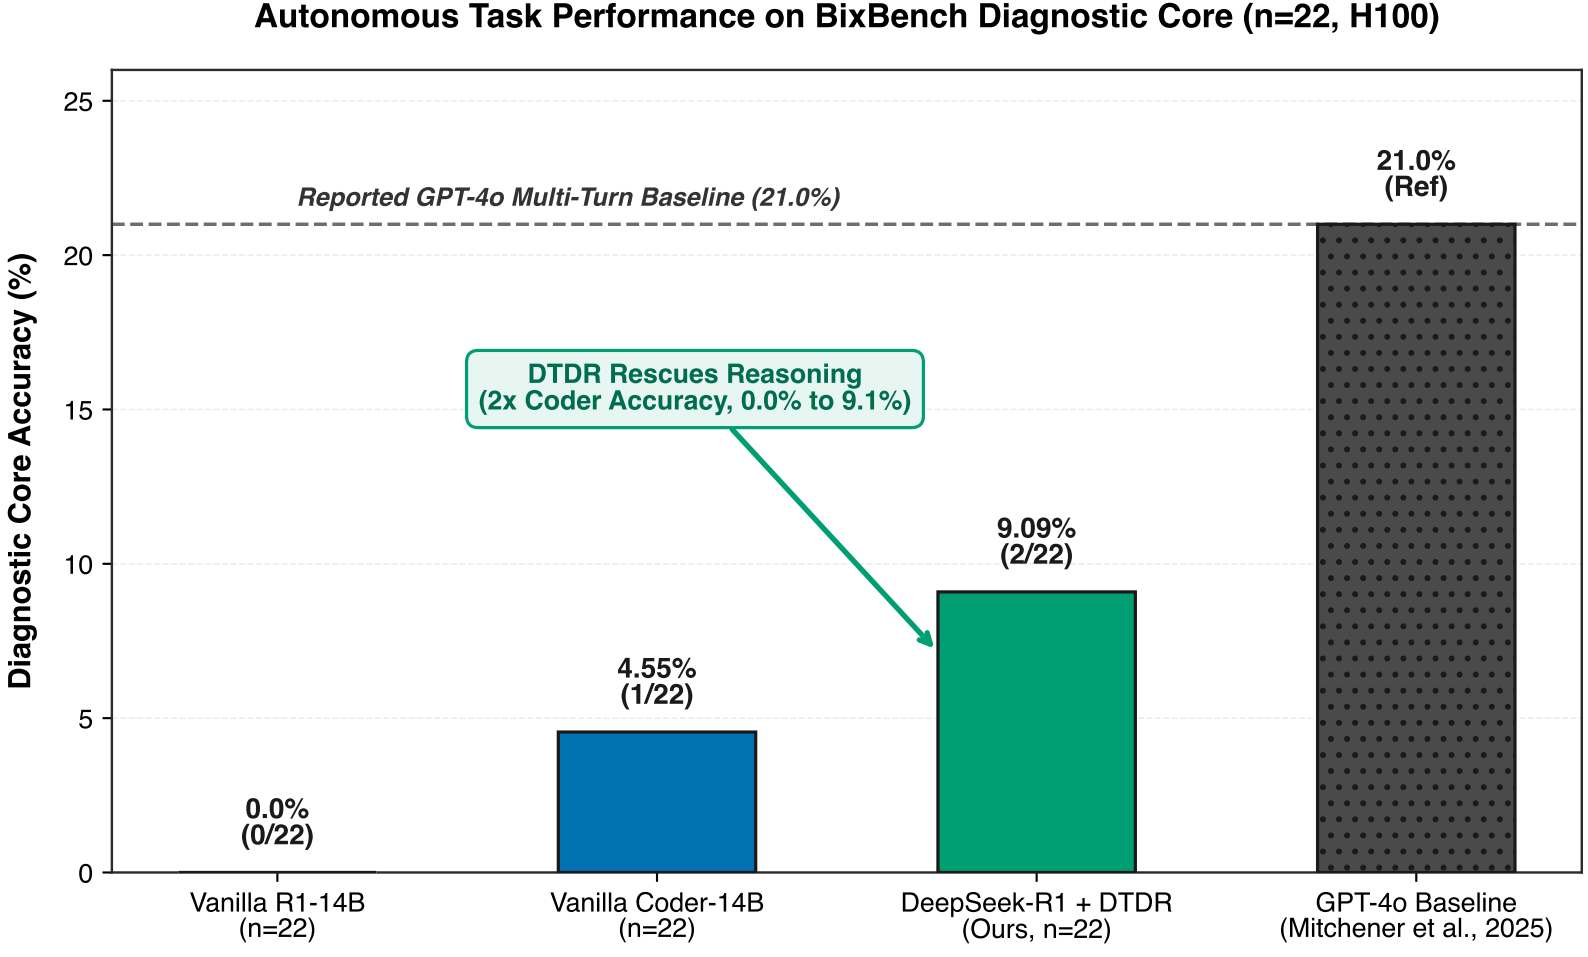
Autonomous Task Performance on BixBench Diagnostic Core (n=22, H100)
Diagnostic Core Accuracy (%)
25
20
15
10
5
0
Reported GPT-4o Multi-Turn Baseline (21.0%)
21.0%
(Ref)
DTDR Rescues Reasoning
(2x Coder Accuracy, 0.0% to 9.1%)
9.09%
(2/22)
4.55%
(1/22)
0.0%
(0/22)
Vanilla R1-14B
(n=22)
Vanilla Coder-14B
(n=22)
DeepSeek-R1 + DTDR
(Ours, n=22)
GPT-4o Baseline
(Mitchener et al., 2025)